

Data Analysis with R — C997

Carl Andrew Perkins

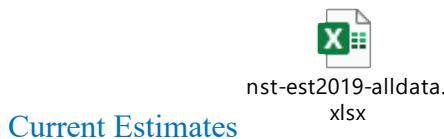
Western Governors University

Data Analysis with R — C997

Below is the process and deliverables by which I imported, scrubbed and created the model in R according to the task's request.

A Create a linear regression analysis with R to predict the size of the population for the state you live in based on the “Current Estimates Data” dataset (see weblink below).

Provide a screenshot of your results.



I have given screenshots below of the linear regression analysis. Note, this was the actual end deliverable and not the beginning of work.

```
#####
# Setting the model parameters
model <- lm(pop~year, data = data)
future <- data.frame(year=c(2020, 2021, 2022, 2023, 2024, 2025))
pred <- predict(model, future)

# view the prediction numbers
pred

# Summarizing descriptive analytics of the model
# (you can see high R value, indicating a strong linear correlation even if you weren't able to look at the plot)
summary(model)

# Combining both population and year data sets for historical and predictive values
total_pop <- c(pop, pred)
total_years <- c(year, 2020, 2021, 2022, 2023, 2024, 2025)
|
# visual Audit of both above
total_pop
total_years

# Plotting data with predictive values
plot(total_years, total_pop, xlab="Year", ylab="Population", main = "Population over Year for Kentucky", type = "b")
```

Console Results

```
Call:
lm(formula = pop ~ year, data = data)

Residuals:
    Min      1Q  Median      3Q     Max 
 -9865  -1413   1938   2694   7403 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) -2.191e+07  1.259e+06 -17.40 1.21e-07 ***
year        1.307e-04  6.251e+02   20.91 2.87e-08 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5678 on 8 degrees of freedom
Multiple R-squared:  0.982,    Adjusted R-squared:  0.9798 
F-statistic: 437.2 on 1 and 8 DF,  p-value: 2.871e-08
```

B. Explain how you prepared the data from part A and how the dataset was imported into R, including screenshots of your results.

The first step was to import libraries like with any R scripts. After the libraries were imported, I proceeded to import the excel data via the readxl library

```
# Make sure I import all the correct libraries and that they are installed
library(tinytex)
library(dplyr)
library(tidyverse)
library(ggplot2)
library(ggthemes)
library(rmarkdown)
library(XLConnect)
library(readxl)
library(devtools)

# Read in excel data source using readxl library
original_est_file <- read_excel("c:\\\\users\\\\andre\\\\Desktop\\\\Extras\\\\WGU\\\\Masters\\\\R for Data Analysts\\\\nst-est2019-alldata.xlsx")
# This original data set includes 57 observations and 151 variables
```

Console Results

```
> # Make sure I import all the correct libraries and that they are installed
> library(tinytex)
> library(dplyr)
> library(tidyverse)
> library(ggplot2)
> library(ggthemes)
> library(rmarkdown)
> library(XLConnect)
> library(readxl)
> library(devtools)
>
>
> # Read in excel data source using readxl library
> original_est_file <- read_excel("c:\\\\users\\\\andre\\\\Desktop\\\\Extras\\\\WGU\\\\Masters\\\\R for Data Analysts\\\\nst-est2019-alldata.xlsx")
```

I then proceeded to wrangle the data by first removing all data but KY and then getting down to only the necessary columns. Please see below

```
# Removing all but the data for Kentucky
df2 <- original_est_file[23,]

#Confirming Kentucky has been pulled (visual audit)
df2[,5]

# Reduces columns to population estimates 2010-2019
df3 <- df2[,1:17]

#Confirming column reductions (visual audit)
df3

# Removes first four columns (unnecessary)
df4 <- df3[,5:17]

#Confirming column reductions (visual audit)
df4

# Creating final data set for analysis (this is only 2010-2019 pop estimates, 1 observation with 10 variables)
data <- df4[-(1:3)]

# One final visual audit
data

# Doing initial summary of each year's observation
summary(data)
```

Console Results

```
> # Removing all but the data for Kentucky
> df2 <- original_est_file[23,]
>
> #Confirming Kentucky has been pulled (visual audit)
> df2
# A tibble: 1 x 151
# ... with 140 more variables: POPESTIMATE2014 <dbl>, POPESTIMATE2015 <dbl>, POPESTIMATE2016 <dbl>, POPESTIMATE2017 <dbl>,
#   ... with 140 more variables: POPESTIMATE2018 <dbl>, POPESTIMATE2019 <dbl>, NPOPCHG_2010 <dbl>, NPOPCHG_2011 <dbl>, NPOPCHG_2012 <dbl>, NPOPCHG_2013 <dbl>,
#   NPOPCHG_2014 <dbl>, NPOPCHG_2015 <dbl>, NPOPCHG_2016 <dbl>, NPOPCHG_2017 <dbl>, NPOPCHG_2018 <dbl>, NPOPCHG_2019 <dbl>,
#   BIRTHS2010 <dbl>, BIRTHS2011 <dbl>, BIRTHS2012 <dbl>, BIRTHS2013 <dbl>, BIRTHS2014 <dbl>, BIRTHS2015 <dbl>, BIRTHS2016 <dbl>,
#   BIRTHS2017 <dbl>, BIRTHS2018 <dbl>, BIRTHS2019 <dbl>, DEATHS2010 <dbl>, DEATHS2011 <dbl>, DEATHS2012 <dbl>, DEATHS2013 <dbl>,
#   DEATHS2014 <dbl>, DEATHS2015 <dbl>, DEATHS2016 <dbl>, DEATHS2017 <dbl>, DEATHS2018 <dbl>, DEATHS2019 <dbl>, NATURALINC2010 <dbl>,
#   NATURALINC2011 <dbl>, NATURALINC2012 <dbl>, NATURALINC2013 <dbl>, NATURALINC2014 <dbl>, NATURALINC2015 <dbl>,
#   NATURALINC2016 <dbl>, NATURALINC2017 <dbl>, NATURALINC2018 <dbl>, NATURALINC2019 <dbl>, INTERNATIONALMIG2010 <dbl>,
#   INTERNATIONALMIG2011 <dbl>, INTERNATIONALMIG2012 <dbl>, INTERNATIONALMIG2013 <dbl>, INTERNATIONALMIG2014 <dbl>,
#   INTERNATIONALMIG2015 <dbl>, INTERNATIONALMIG2016 <dbl>, INTERNATIONALMIG2017 <dbl>, INTERNATIONALMIG2018 <dbl>,
#   INTERNATIONALMIG2019 <dbl>, DOMESTICMIG2010 <dbl>, DOMESTICMIG2011 <dbl>, DOMESTICMIG2012 <dbl>, DOMESTICMIG2013 <dbl>,
#   DOMESTICMIG2014 <dbl>, DOMESTICMIG2015 <dbl>, DOMESTICMIG2016 <dbl>, DOMESTICMIG2017 <dbl>, DOMESTICMIG2018 <dbl>,
#   DOMESTICMIG2019 <dbl>, NETMIG2010 <dbl>, NETMIG2011 <dbl>, NETMIG2012 <dbl>, NETMIG2013 <dbl>, NETMIG2014 <dbl>, NETMIG2015 <dbl>,
#   NETMIG2016 <dbl>, NETMIG2017 <dbl>, NETMIG2018 <dbl>, NETMIG2019 <dbl>, RESIDUAL2010 <dbl>, RESIDUAL2011 <dbl>,
#   RESIDUAL2012 <dbl>, RESIDUAL2013 <dbl>, RESIDUAL2014 <dbl>, RESIDUAL2015 <dbl>, RESIDUAL2016 <dbl>, RESIDUAL2017 <dbl>,
#   RESIDUAL2018 <dbl>, RESIDUAL2019 <dbl>, RBIRTH2011 <dbl>, RBIRTH2012 <dbl>, RBIRTH2013 <dbl>, RBIRTH2014 <dbl>, RBIRTH2015 <dbl>,
#   RBIRTH2016 <dbl>, RBIRTH2017 <dbl>, RBIRTH2018 <dbl>, RBIRTH2019 <dbl>, RDEATH2011 <dbl>, RDEATH2012 <dbl>, RDEATH2013 <dbl>,
#   RDEATH2014 <dbl>, RDEATH2015 <dbl>, ...
> df2[,5]
# A tibble: 1 x 1
# ... with 140 more variables: POPESTIMATE2010 <dbl>, POPESTIMATE2011 <dbl>, POPESTIMATE2012 <dbl>, POPESTIMATE2013 <dbl>,
#   ... with 140 more variables: POPESTIMATE2014 <dbl>, POPESTIMATE2015 <dbl>, POPESTIMATE2016 <dbl>, POPESTIMATE2017 <dbl>,
#   POPESTIMATE2018 <dbl>, POPESTIMATE2019 <dbl>
>
> # Confirming column reductions (visual audit)
> df3
# A tibble: 1 x 17
# ... with 140 more variables: POPESTIMATE2010 <dbl>, POPESTIMATE2011 <dbl>, POPESTIMATE2012 <dbl>, POPESTIMATE2013 <dbl>,
#   ... with 140 more variables: POPESTIMATE2014 <dbl>, POPESTIMATE2015 <dbl>, POPESTIMATE2016 <dbl>, POPESTIMATE2017 <dbl>,
#   POPESTIMATE2018 <dbl>, POPESTIMATE2019 <dbl>
>
> # Removes first four columns (unnecessary)
> df4 <- df3[,5:17]
>
> #Confirming column reductions (visual audit)
> df4
# A tibble: 1 x 13
# ... with 140 more variables: POPESTIMATE2010 <dbl>, POPESTIMATE2011 <dbl>, POPESTIMATE2012 <dbl>, POPESTIMATE2013 <dbl>,
#   ... with 140 more variables: POPESTIMATE2014 <dbl>, POPESTIMATE2015 <dbl>, POPESTIMATE2016 <dbl>, POPESTIMATE2017 <dbl>,
#   POPESTIMATE2018 <dbl>, POPESTIMATE2019 <dbl>
>
> # Creating final data set for analysis (this is only 2010-2019 pop estimates, 1 observation with 10 variables)
> data <- df4[,-(1:3)]
>
> # One final visual audit
> data
# A tibble: 1 x 10
# ... with 140 more variables: POPESTIMATE2010 <dbl>, POPESTIMATE2011 <dbl>, POPESTIMATE2012 <dbl>, POPESTIMATE2013 <dbl>,
#   ... with 140 more variables: POPESTIMATE2014 <dbl>, POPESTIMATE2015 <dbl>, POPESTIMATE2016 <dbl>, POPESTIMATE2017 <dbl>,
#   POPESTIMATE2018 <dbl>, POPESTIMATE2019 <dbl>
>
> # Doing initial summary of each year's observation
> summary(data)
POPESTIMATE2010    POPESTIMATE2011    POPESTIMATE2012    POPESTIMATE2013    POPESTIMATE2014    POPESTIMATE2015    POPESTIMATE2016
Min. :4348181     Min. :4339367     Min. :4339333     Min. :4348181     Min. :4369821     Min. :4386346     Min. :4404659
1st Qu.:4348181   1st Qu.:4339367   1st Qu.:4339333   1st Qu.:4348181   1st Qu.:4369821   1st Qu.:4386346   1st Qu.:4404659
Median:4348181   Median:4339367   Median:4339333   Median:4348181   Median:4369821   Median:4386346   Median:4404659
Mean :4348181    Mean :4339367    Mean :4339333    Mean :4348181    Mean :4369821    Mean :4386346    Mean :4404659
3rd Qu.:4348181  3rd Qu.:4339367  3rd Qu.:4339333  3rd Qu.:4348181  3rd Qu.:4369821  3rd Qu.:4386346  3rd Qu.:4404659
Max. :4348181   Max. :4339367   Max. :4339333   Max. :4348181   Max. :4369821   Max. :4386346   Max. :4404659
POPESTIMATE2017    POPESTIMATE2018    POPESTIMATE2019
Min. :4452268     Min. :4461153     Min. :4461153
1st Qu.:4452268   1st Qu.:4461153   1st Qu.:4461153
Median:4452268   Median:4461153   Median:4461153
Mean :4452268    Mean :4461153    Mean :4461153
3rd Qu.:4452268  3rd Qu.:4461153  3rd Qu.:4461153
Max. :4452268   Max. :4461153   Max. :4461153

```

C. Create an R script that will tabulate a statistical description of the model using R's summary() function and provide a screenshot of your results.

I ran an initial summary and you can see those statistics on the original data set above.

However, I also ran a summary on the model itself, please see below.

```
#####
# Setting the model parameters
model <- lm(pop~year, data = data)
future <- data.frame(year=c(2020, 2021, 2022, 2023, 2024, 2025))
pred <- predict(model, future)

# View the prediction numbers
pred

# Summarizing descriptive analytics of the model
# (you can see high R value, indicating a strong linear correlation even if you weren't able to look at the plot)
summary(model)
```

Console Results

```
> # Summarizing descriptive analytics of the model
> # (you can see high R value, indicating a strong linear correlation even if you weren't able to look at the plot)
> summary(model)

Call:
lm(formula = pop ~ year, data = data)

Residuals:
    Min     1Q Median     3Q    Max 
-9865 -1413  1938  2694  7403 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) -2.191e+07  1.259e+06 -17.40 1.21e-07 ***
year         1.307e+04  6.251e+02   20.91 2.87e-08 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5678 on 8 degrees of freedom
Multiple R-squared:  0.982,    Adjusted R-squared:  0.9798 
F-statistic: 437.2 on 1 and 8 DF,  p-value: 2.871e-08
```

D. Predict the population size of your state in five years using a linear regression

from part A and provide a screenshot of your results.

This follow is the original plotting of the data all the way through the modeling and plotting of the regression prediction.

```
#####
# This is Population over year for the state of Kentucky

# Creating population values based upon data frame
pop <- c(4348181, 4369821, 4386346, 4404659, 4414349, 4425976, 4438182, 4452268, 4461153, 4467673)

# visual audit of population
pop

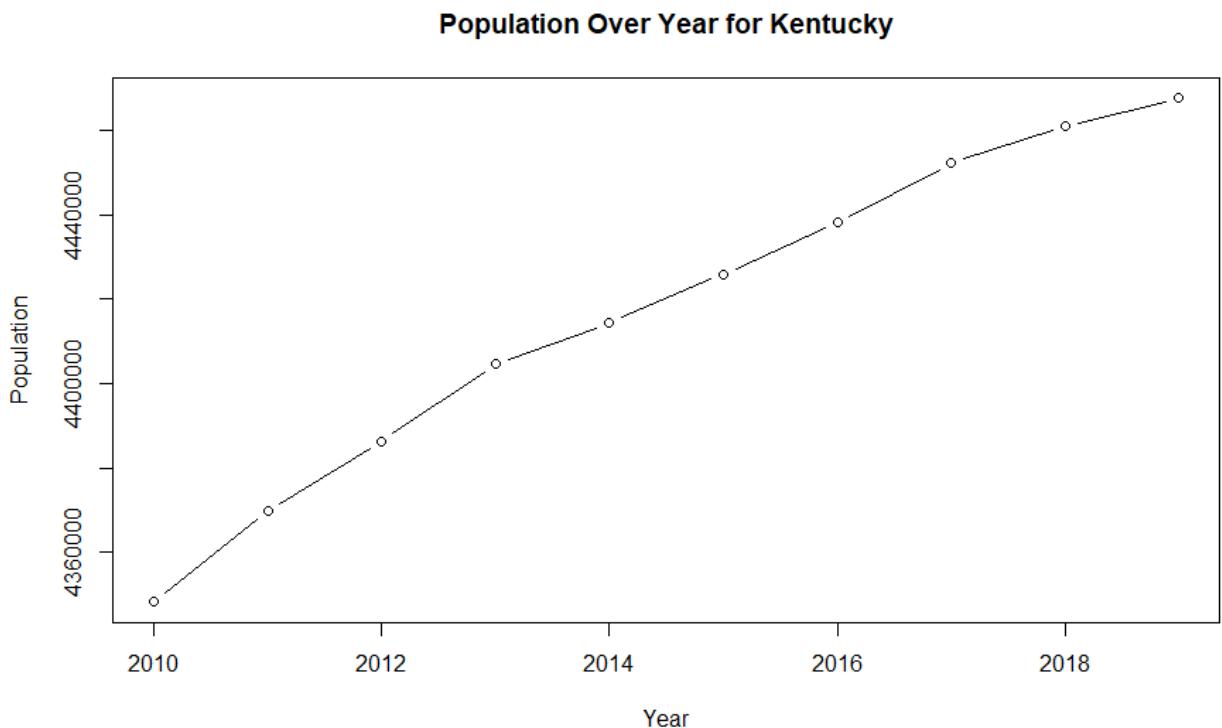
# Creating year values based upon data frame
year <- c(2010, 2011, 2012, 2013, 2014, 2015, 2016, 2017, 2018, 2019)

# visual audit of year
year

#Actually creating the plot using the previously created values
plot(year, pop,xlab="year",ylab="Population", main = "Population over year for Kentucky", type = "b")
```

Console and Plot Results for Historical

```
> # This is Population over year for the state of Kentucky
>
> # creating population values based upon data frame
> pop <- c(4348181, 4369821, 4386346, 4404659, 4414349, 4425976, 4438182, 4452268, 4461153, 4467673)
>
> # visual audit of population
> pop
[1] 4348181 4369821 4386346 4404659 4414349 4425976 4438182 4452268 4461153 4467673
>
> # Creating year values based upon data frame
> year <- c(2010, 2011, 2012, 2013, 2014, 2015, 2016, 2017, 2018, 2019)
>
> # visual audit of year
> year
[1] 2010 2011 2012 2013 2014 2015 2016 2017 2018 2019
>
> #Actually creating the plot using the previously created values
> plot(year, pop,xlab="Year",ylab="Population", main = "Population Over Year for Kentucky", type = "b")
```



R code for predictive model

```
#####
# Setting the model parameters
model <- lm(pop~year, data = data)
future <- data.frame(year=c(2020, 2021, 2022, 2023, 2024, 2025))
pred <- predict(model, future)

# view the prediction numbers
pred

# Summarizing descriptive analytics of the model
# (you can see high R value, indicating a strong linear correlation even if you weren't able to look at the plot)
summary(model)

# Combining both population and year data sets for historical and predictive values
total_pop <- c(pop, pred)
total_years <- c(year, 2020, 2021, 2022, 2023, 2024, 2025)

# visual Audit of both above
total_pop
total_years

# Plotting data with predictive values
plot(total_years, total_pop, xlab="Year",ylab="Population", main = "Population over Year for Kentucky", type = "b")
```

Console and Plot Results in conjunction with predictive model

```

> ##### MODEL #####
> # Setting the model parameters
> model <- lm(pop~year, data = data)
> future <- data.frame(year=c(2020, 2021, 2022, 2023, 2024, 2025))
> pred <- predict(model, future)
>
> # view the prediction numbers
> pred
   1     2     3     4     5     6
4488746 4501816 4514886 4527956 4541026 4554096
>
> # Summarizing descriptive analytics of the model
> # (you can see high R value, indicating a strong linear correlation even if you weren't able to look at the plot)
> summary(model)

Call:
lm(formula = pop ~ year, data = data)

Residuals:
    Min      1Q Median      3Q      Max  
-9865  -1413   1938   2694   7403  

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) -2.191e+07 1.259e+06 -17.40 1.21e-07 ***
year        1.307e+04 6.251e+02   20.91 2.87e-08 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5678 on 8 degrees of freedom
Multiple R-squared:  0.982, Adjusted R-squared:  0.9798 
F-statistic: 437.2 on 1 and 8 DF,  p-value: 2.871e-08

> # This is Population over year for the state of Kentucky
>
> # Creating population values based upon data frame
> pop <- c(4348181, 4369821, 4386346, 4404659, 4414349, 4425976, 4438182, 4452268, 4461153, 4467673)
>
> # visual audit of population
> pop
[1] 4348181 4369821 4386346 4404659 4414349 4425976 4438182 4452268 4461153 4467673

> # Creating year values based upon data frame
> year <- c(2010, 2011, 2012, 2013, 2014, 2015, 2016, 2017, 2018, 2019)
>
> # visual audit of year
> year
[1] 2010 2011 2012 2013 2014 2015 2016 2017 2018 2019
>
> #Actually creating the plot using the previously created values
> plot(year, pop,xlab="Year",ylab="Population", main = "Population over Year for Kentucky", type = "b")
> # Setting the model parameters
> model <- lm(pop~year, data = data)
> future <- data.frame(year=c(2020, 2021, 2022, 2023, 2024, 2025))
> pred <- predict(model, future)
>
> # view the prediction numbers
> pred
   1     2     3     4     5     6
4488746 4501816 4514886 4527956 4541026 4554096
>
> # summarizing descriptive analytics of the model
> # (you can see high R value, indicating a strong linear correlation even if you weren't able to look at the plot)
> summary(model)

Call:
lm(formula = pop ~ year, data = data)

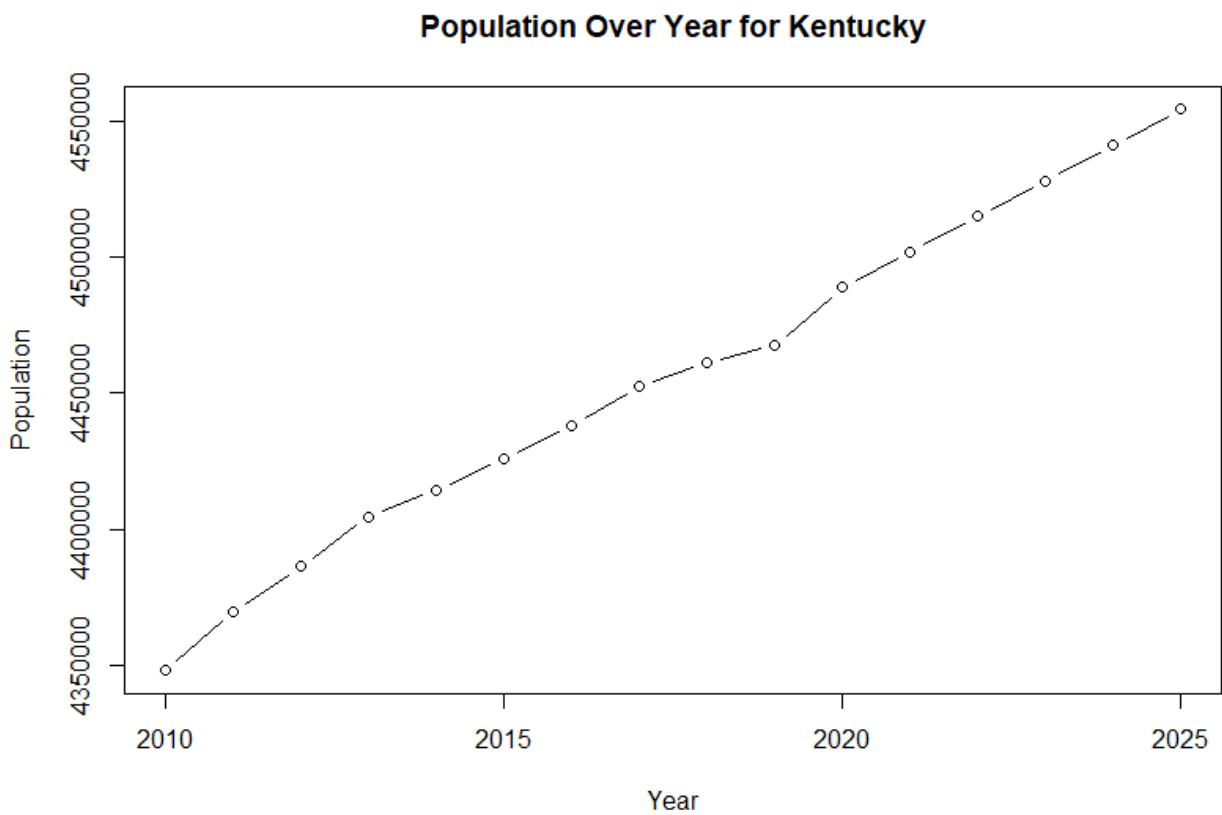
Residuals:
    Min      1Q Median      3Q      Max  
-9865  -1413   1938   2694   7403  

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) -2.191e+07 1.259e+06 -17.40 1.21e-07 ***
year        1.307e+04 6.251e+02   20.91 2.87e-08 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5678 on 8 degrees of freedom
Multiple R-squared:  0.982, Adjusted R-squared:  0.9798 
F-statistic: 437.2 on 1 and 8 DF,  p-value: 2.871e-08

> # Combining both population and year data sets for historical and predictive values
> total_pop <- c(pop, pred)
> total_years <- c(year, 2020, 2021, 2022, 2023, 2024, 2025)
>
> # Visual Audit of both above
> total_pop
   1     2     3     4     5
4348181 4369821 4386346 4404659 4414349 4425976 4438182 4452268 4461153 4467673 4488746 4501816 4514886 4527956 4541026
   6
4554096
> total_years
[1] 2010 2011 2012 2013 2014 2015 2016 2017 2018 2019 2020 2021 2022 2023 2024 2025
>
> # Plotting data with predictive values
> plot(total_years, total_pop, xlab="Year",ylab="Population", main = "Population Over Year for Kentucky", type = "b")

```



References

How to Create a Data Frame from Scratch in R. (2019, November 21). Dummies.

Retrieved from: <https://www.dummies.com/programming/r/how-to-create-a-data-frame-from-scratch-in-r/>